



KV allocation policy comparison (OpenVLA, Azure conv, n=2000, scale=3.0, SLO=500ms — moderate-load regime)


(a) Throughput



(b) E2E latency p99



(c) Admitted requests (aggression)



(d) Tokens recomputed (wasted work)

